

NITESH RANJAN SINGH

Glasgow, UK | +44 7900 945064 | niteshranjansingh85389@gmail.com | [linkedin.com/in/nitesh0007](https://www.linkedin.com/in/nitesh0007) | [GitHub](#) | [Portfolio](#)

Work authorisation: UK Graduate Route visa eligible.

PROFESSIONAL SUMMARY

Data & AI Engineer with **3+ years** building regulated cloud data pipelines for pharmaceutical clients. Azure, Databricks, Python and SQL across **10M+ records per batch**, complemented by deployed agentic AI projects with testing and observability. MSc Data Science, Glasgow; Azure, Fabric and Databricks certified.

TECHNICAL SKILLS

- **Data Engineering:** Python · SQL · PySpark · ETL/ELT · batch pipelines · CDC · schema validation
- **Azure & Data Platforms:** Azure Data Factory · Databricks · ADLS Gen2 · Synapse · Delta Lake
- **Cloud Infrastructure:** AWS (Lambda, S3) · Snowflake · AKS · Docker · GitHub Actions · Git · CI/CD
- **Data Modelling:** Dimensional/star schema · SCD · medallion architecture · source-to-target mapping
- **Quality & Governance:** Data quality · data lineage · metadata · monitoring · observability
- **GenAI & Agents:** LangGraph · MCP · RAG · FastAPI · Pydantic · human-in-the-loop workflows
- **Client Delivery:** Stakeholder requirements · documentation · Claude/Copilot AI-assisted development

PROFESSIONAL EXPERIENCE

Data System & CRM Migration Intern

Jun 2026 – Present

Street Connect · Glasgow, UK · alongside MSc

- Support CRM migration through Python load routines, source-to-target mappings and reconciliation checks; maintain reporting definitions ahead of cutover.

Data Engineer

Dec 2022 – Aug 2025

IQVIA | GSK Omnichannel Data Platform · Gurugram, India · promoted from Data Analyst

- Built Azure Databricks/PySpark CDC pipelines for 10M+ records per batch across UK, Germany and Italy; schema-validation gates achieved 99.8% accuracy and partition tuning cut processing time by 35%.
- Metadata-driven Azure Data Factory ETL with SQL quality logging reduced manual QA by 65% and new-market onboarding from 3 weeks to 1 week.
- A reusable PySpark framework checked nulls, duplicates, reconciliation and drift, improving reliability by 40% and reducing QA time by 60%; adopted by 2 other client teams.
- Dimensional models and Synapse reporting datasets supported 50+ Power BI KPIs for 200+ users, turning complex client requirements into consistent analytical outputs.
- Documented schema contracts, validation severity and transformation behaviour so analysts could maintain pipelines and interpret quality failures.

Analyst Trainee

Jan 2022 – Dec 2022

Axtria | Johnson & Johnson MDM and AstraZeneca O2 · Noida, India

- Developed Azure Data Factory, Databricks and ADLS Gen2 pipelines for healthcare analytics; distributed PySpark processing improved speed by 54%.

KEY PROJECTS

DataSentinel: Data Reliability Engineer | Live FastAPI/pandas/Pydantic application tracing pipeline failures to Git changes.

- Deterministic data quality checks caught a demo job silently dropping 60.2% of rows despite reporting success; Git diff evidence identified the regressing file, line and commit.
- Human-in-the-loop approval, reverse patches, rerun/re-profile gates and rollback control remediation; optional evidence-grounded LLM explanations and 62 backend tests.

LumaLearn SmartReco | Live agentic recommender: FastAPI, LangGraph, Chroma and GPT-4o-mini.

- A bounded analyse/retrieve/grade/refine/generate workflow validates generated IDs against retrieved candidates; typed outputs and SQL/vector outbox reconciliation protect consistency.
- Per-node latency/token traces, cost gates and Docker deployment support observability; 10 automated tests and 2 passing GitHub Actions workflows provide CI/CD quality gates.

EDUCATION

MSc Data Science

Sep 2025 – Sep 2026

University of Glasgow, UK | £10,000 Excellence Scholarship

Dissertation: Provider Fairness in Multi-Stakeholder Recommendation, combining graph models with grounded LLM explanations.

B.Tech Computer Science (AI & ML)

Jun 2018 – May 2022

UPES, India

CERTIFICATIONS & AWARDS

Certifications: Azure Data Engineer Associate (DP-203) | Microsoft Fabric Data Engineer Associate (DP-700) | Databricks Data Engineer Associate | IBM Data Science Professional

Awards: Bronze Medal, Kaggriculture (Kaggle) | 2nd Place & \$3,000, Kaggle Agentic Legal IR | Runner-Up, GUTS 2025 Hackathon | IQVIA Impact Award FY23-24